

1. Socket programming is a programming paradigm that allows communication between two computer processes over a network. It provides a way to establish a connection between a client and a server and enables them to exchange data.

2. In the context of networking, a client is a program or device that requests services or resources from another program or device called the server. The client initiates a connection and sends requests to the server, which responds with the requested data or performs the requested action.

3. Inter Process Communication (IPC) refers to the mechanisms and techniques used by various processes or programs to communicate and share data with each other within a computer system. It enables processes to coordinate their activities and exchange information.

4. There are several IPC mechanisms, including:

- Sockets: Communication over a network using TCP/IP or UDP protocols.
- Pipes: A unidirectional communication channel between two related processes.
- Shared Memory: Processes share a common memory region for communication.
- Message Passing: Processes send and receive messages through a communication channel.
- Remote Procedure Calls (RPC): Processes can call functions or procedures in a remote process as if they were local.

5. A socket is an endpoint for communication between two machines over a network. It represents a combination of IP address and port number. Sockets can be used to establish connections, send and receive data, and close connections.

6. TCP (Transmission Control Protocol) and UDP (User Datagram Protocol) are transport layer protocols used in socket communication. The main differences between them are:

- TCP provides reliable, ordered, and error-checked delivery of data, while UDP is unreliable and does not guarantee delivery.
- TCP establishes a connection before sending data, while UDP does not require a connection.
- TCP is stream-oriented, which means it sends a continuous stream of data, while UDP is message-oriented, sending individual packets.
- TCP performs congestion control to ensure reliable delivery, while UDP does not have built-in congestion control.

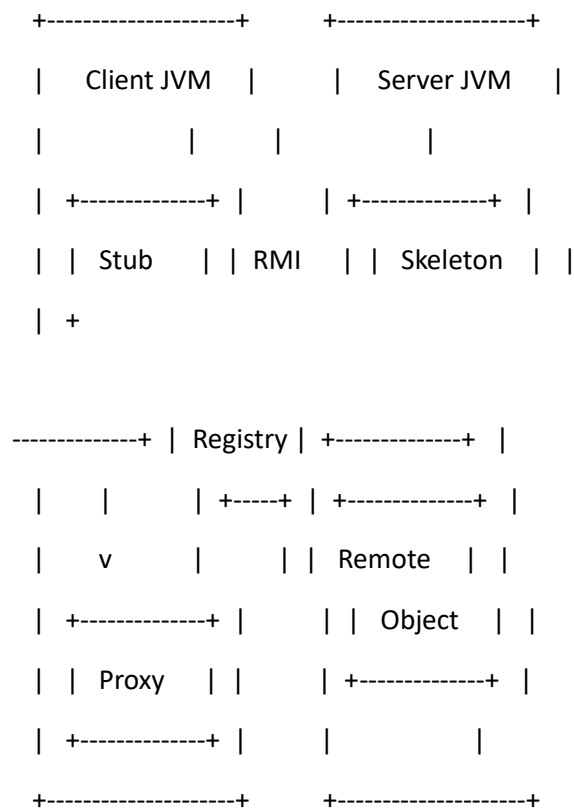
7. Shared memory programming is a technique in which multiple processes can access and manipulate a shared region of memory concurrently. It allows processes to communicate and share data by reading from and writing to the same memory area.

8. In the context of networking, a port is a numerical identifier used to distinguish between different network services running on a single device. It helps in directing incoming network traffic to the appropriate service or application. Applications use specific port numbers to communicate with other applications on remote machines. For example, HTTP typically uses port 80, while HTTPS uses port 443.

9. Heterogeneity refers to the concept of systems or networks consisting of diverse elements that may have different architectures, platforms, or data formats. In a heterogeneous environment, different components or systems need to interact and communicate effectively despite their differences.

10. Marshalling and unmarshalling are techniques used for the serialization and deserialization of data when transmitting it between different systems or languages. Marshalling involves converting data objects into a format suitable for transmission, typically a byte stream. Unmarshalling is the reverse process of reconstructing data objects from the transmitted format.

11. RMI (Remote Method Invocation) is a Java technology that allows the invocation of methods on remote objects. It enables distributed computing in Java by providing a mechanism for objects in one JVM (Java Virtual Machine) to invoke methods on objects residing in another JVM. Here is a simplified diagram illustrating the components involved in RMI:



12. Binding in the context of RMI refers to the process of associating a remote object with a name or address that clients can use to locate and access the object. It involves registering the object with the RMI registry, which acts as a central directory for remote objects.

13. The RMI registry serves as a central registry or naming service in RMI. It allows clients to look up remote objects by their registered names and obtain references to those objects. The RMI registry needs to be started first so that clients can discover and access the remote objects it hosts.

14. In RMI, "UnicastRemoteObject" is a class that provides a convenient way to create and export remote objects. "lookup()" is a method used to look up a remote object in the RMI registry based on its registered name. "rebind()" is a method used to bind or rebind a remote object to a name in the RMI registry.

15. In RMI, a stub represents the client-side proxy for the remote object. It acts as a local representative of the remote object and handles the communication details on behalf of the client. A skeleton, on the other hand, resides on the server-side and receives incoming requests from the client, dispatching them to the appropriate remote object.

16. An Exception is a general term referring to an abnormal condition or error that occurs during the execution of a program. In the context of RMI, a Remote Exception specifically indicates that an exception occurred during a remote method invocation or communication with a remote object.

17. CORBA (Common Object Request Broker Architecture) is a middleware technology that enables communication and interaction between distributed objects across different platforms and languages. It provides a standardized approach for building distributed systems by defining interfaces, communication protocols, and object services.

18. The CORBA architecture consists of several components, including:

- Object Request Broker (ORB): The central component responsible for locating and directing method invocations to the appropriate objects.
- Interface Definition Language (IDL): A language used to define the interfaces of CORBA objects.
- Object Services: Predefined services that can be utilized by CORBA objects, such as naming, event handling, and persistence.
- Object Adapters: Bridges between CORBA objects and the ORB, handling object lifecycle and invocation details.

19. CORBA works by allowing distributed objects to communicate through the Object Request Broker (ORB). Clients make requests to the ORB, which locates the target objects based on their registered names and routes the requests to the appropriate servers. The servers execute the requested methods and return the results back to the clients through the ORB.

20. CORBA can support both synchronous and asynchronous application models. In a synchronous application, the client blocks until it receives a response from the server. In an asynchronous application, the client can continue its execution without waiting for a response, and the server sends the response at a later time.

21. ORB (Object Request Broker) is a middleware component responsible for managing the communication between distributed objects in CORBA. It provides services such as object location, method invocation, and parameter marshalling, allowing objects to interact seamlessly across different platforms and languages.

22. IDL (Interface Definition Language) is a language used to define the

interfaces of CORBA objects. It provides a platform-independent way to describe the methods, data structures, and exceptions that can be used by objects in a distributed system. IDL allows clients and servers to understand and communicate with each other regardless of their programming languages.

23. Object Request Broker Daemon (ORBD) is a separate process or service that runs in the background and provides the ORB functionality in CORBA. It acts as the central coordination point for communication between distributed objects and manages the lifecycle of objects and their interactions.

24. Middleware is software that sits between the operating system and the applications, providing a set of services and abstractions to facilitate the development of distributed systems. It hides the complexities of network communication, concurrency, and platform differences, allowing developers to focus on application logic.

25. Examples of middleware include CORBA, Java RMI, Message Queueing Middleware (e.g., RabbitMQ), Web Services (e.g., SOAP, REST), Enterprise Service Bus (ESB), and Distributed Object Middleware (e.g., DCOM).

26. Middleware has various uses, including:

- Enabling communication and interaction between distributed components or systems.
- Providing mechanisms for data sharing and synchronization.
- Facilitating interoperability between different platforms, languages, and protocols.
- Supporting the development of scalable and distributed applications.
- Providing services such as security, transactions, and persistence in a distributed environment.

27. CORBA finds applications in various domains, including telecommunications, finance, healthcare, and enterprise systems. Some of the applications of CORBA include:

- Distributed telecommunications systems for call control and service management.
- Enterprise applications with distributed components and services.
- Complex software systems requiring interoperability between different platforms and languages.
- Real-time systems requiring efficient and reliable communication.

28. The main differences between RMI and CORBA are:

- RMI is specific to the Java language, while CORBA supports multiple programming languages.
- RMI uses Java serialization for object communication, while CORBA uses a standardized Interface Definition Language (IDL) and an ORB.
- RMI is typically used in Java-centric environments, while CORBA is more widely adopted in heterogeneous and enterprise systems.

29. RMI is implemented only in Java because it leverages Java's built-in features and capabilities, such as Java object serialization, dynamic class loading, and garbage collection. It tightly integrates with the Java language and runtime environment, making it easier to develop and deploy distributed Java applications.

30. MPI (Message Passing Interface) is a standard programming model and library for message-passing parallel computing. It provides a portable and efficient way to develop parallel programs that can run on distributed memory systems, such as clusters or supercomputers.

31. MPI is commonly used in high-performance computing (HPC) applications where parallel processing and efficient communication between multiple processes are required. It is used in scientific simulations, numerical computations, weather modeling, molecular dynamics, and other computationally intensive tasks.

32. The rank in MPI is a unique identifier assigned to each process participating in a parallel program. It helps in distinguishing and addressing different processes in a collective communication or coordination.

33. MPI provides various operations for communication and coordination among processes, including point-to-point communication (send, receive), collective operations (broadcast, reduce, scatter, gather), and synchronization primitives (barrier, synchronization).

34. MPI supports various data types, including basic types such as integers, floating-point numbers, and characters, as well as derived types such as arrays, structures, and user-defined data structures. These data types enable efficient and flexible data exchange between processes.

35. The MPI architecture consists of multiple processes running on different compute nodes or machines. Each process has its own memory space and executes the program logic. The processes communicate and coordinate with

each other using MPI function calls, passing messages and data between them.

36. `MPI_ABORT` is a function call in MPI that aborts the execution of all processes within a communicator. It is typically used to terminate the program in case of an unrecoverable error or exceptional condition.

37. `MPI_FINALIZE` is a function call in MPI that marks the end of the MPI program. It is called after the completion of all MPI communication and before the program exits. `MPI_FINALIZE` releases any resources allocated by the MPI library and performs necessary cleanup.

38. MPI_ABORT and MPI_FINALIZE serve different purposes:

- MPI_ABORT is used to forcefully terminate the program and all processes when an unrecoverable error occurs. It does not ensure proper cleanup or communication completion.
- MPI_FINALIZE is called at the end of the program to gracefully terminate the MPI execution. It allows processes to perform necessary cleanup, complete ongoing communication, and release allocated resources.

39. A logical clock is a mechanism used in distributed systems to order events that occur in different processes. It provides a logical ordering based on causality, even when physical clocks in different processes are not synchronized.

40. It is necessary to synchronize clocks in distributed real-time systems to ensure consistent and coordinated behavior. Synchronized clocks enable accurate timestamping of events, coordination of distributed activities, and synchronization of time-sensitive operations.

41. The Berkeley algorithm is a clock synchronization algorithm used in distributed systems. It works by periodically exchanging time information between processes and adjusting the local clock based on the received information. The algorithm aims to minimize the clock skew between processes.

42. Other algorithms for clock synchronization in distributed systems include the Network Time Protocol (NTP), Precision Time Protocol (PTP), and Cristian's algorithm. These algorithms utilize different techniques, such as clock adjustment, message exchange, and timestamp comparison, to achieve clock synchronization.

43. Race condition refers to a situation in concurrent programming where the outcome of the program depends on the relative timing or interleaving of multiple threads or processes. It occurs when two or more threads/processes access shared resources or variables concurrently, and the final result depends on the order of execution.

44. Deadlock is a situation in concurrent programming where two or more processes are unable to proceed because each is waiting for a resource held by the other. It leads to a permanent blocking of processes and can result in the system becoming unresponsive.

45. Starvation refers to a situation where a process or thread is unable to make progress because it is continually denied access to a required resource. It occurs when resources are allocated to other processes indefinitely, leaving a particular process waiting indefinitely.

46. Mutual Exclusion is a property that ensures that only one process or thread can access a shared resource or variable at any given time. It prevents concurrent access that could lead to race conditions or data inconsistency.

47. Token ring algorithm is a distributed algorithm used for achieving mutual exclusion in a distributed system. It involves passing a token from one process to another in a ring-like fashion. Only the process holding the token has the permission to access the shared resource, ensuring mutual exclusion.

48. Mutual exclusion can be achieved using token-based algorithms. In such algorithms, a token circulates among the processes in a distributed system. Only the process holding the token can enter the critical section and access the shared resource. Other processes must wait until they receive the token.

49. In distributed systems, a process coordinator is a designated process responsible for managing and coordinating certain tasks or activities among multiple processes. The process coordinator often takes on additional responsibilities such as election coordination, resource allocation, or decision-making in distributed algorithms.

50. The responsibilities of a process coordinator may include:

- Coordination and synchronization of distributed activities.
- Managing resources and allocating them to processes.
- Handling communication and message passing between processes
- Coordinating distributed algorithms, such as leader election or consensus protocols.
- Monitoring and fault detection in the distributed system.

51. Election algorithms are used in distributed systems to elect a leader or coordinator among a group of processes. The need for election arises when a process coordinator fails or when a new coordinator needs to be selected due to system changes or failures.

52. Centralized algorithms rely on a central authority or coordinator to manage the election process. The central authority receives election requests and determines the new leader. Decentralized algorithms distribute the election responsibility among multiple processes, and the processes collaboratively elect a new leader.

53. The working of an election algorithm depends on the specific algorithm used. In the Ring algorithm, processes are organized in a logical ring structure, and the election message circulates through the ring until it reaches the highest-ranked process. In the Bully algorithm, processes with higher ranks challenge lower-ranked processes, and the highest-ranked process becomes the leader.

54. The main differences between the Bully and Ring algorithms for leader election are:

- The Bully algorithm assumes a hierarchy of processes based on their ranks, while the Ring algorithm forms a logical ring structure.
- The Bully algorithm involves higher-ranked processes challenging lower-ranked processes, while the Ring algorithm passes an election message around the ring.
- The Bully algorithm ensures that the highest-ranked process becomes the leader, while the Ring algorithm selects the process with the highest rank.

55. In the context of distributed systems, a token is a special message or token object that is passed between processes to control access to a shared resource. The process holding the token has the permission to access the shared resource, ensuring exclusive access and preventing concurrent access conflicts.

56. A web service is a software system designed to enable interoperable communication between different applications or systems over a network. It provides a standardized way to expose and consume functionality, allowing applications to interact and share data in a platform-independent manner.

57. The architecture of web services typically involves the following components:

- Provider: The entity or application that publishes the web service and makes its functionality available to other applications.
- Requestor: The entity or application that consumes or uses the web service by sending requests and receiving responses.
- Service registry: A directory or repository that stores information about available web services, including their locations and interfaces.
- Broker: A component or intermediary that facilitates the discovery, binding, and coordination between providers and requestors.

58. WSDL (Web Services Description Language) is an XML-based language used to describe the interface and functionality of a web service. It defines the operations, input/output messages, data types, and protocols used by the web service.

59. There are different types of web services, including:

- SOAP (Simple Object Access Protocol): Web services that use the SOAP protocol for communication. They typically rely on XML-based messages and are described using WSDL.
- REST (Representational State Transfer): Web services that follow the principles of REST architectural style. They use standard HTTP methods (GET, POST, PUT, DELETE) and operate on resources identified by URLs.
- JSON-RPC: Web services that use the JSON (JavaScript Object Notation) format for data serialization and follow the RPC (Remote Procedure Call) paradigm.

60. SOAP (Simple Object Access Protocol) and REST (Representational State Transfer) are two different approaches to building web services:

- SOAP is a protocol that uses XML-based messages and supports more complex operations and data types. It provides a standardized and extensible way for communication between systems but can be more verbose and require additional processing overhead.
- REST is an architectural style that uses standard HTTP methods and operates on resources identified by URLs. It provides a lightweight and straightforward approach to building web services

but may lack some of the advanced features provided by SOAP.

61. Examples of web services include:

- Weather service: A web service that provides weather information based on location and date/time parameters.
- Payment gateway: A web service that handles online payment processing for e-commerce websites.
- Geocoding service: A web service that converts addresses into geographic coordinates and vice versa.
- Social media API: Web services provided by social media platforms to interact with user profiles, posts, and social graphs.

62. Web services find applications in various domains, including:

- E-commerce: Web services enable secure online transactions, payment processing, and integration with third-party services.
- Enterprise integration: Web services facilitate communication and data exchange between different systems and applications within an organization.
- Mobile applications: Web services provide APIs for mobile apps to access remote resources, such as data storage, social media integration, or location services.
- Cloud computing: Web services play a crucial role in cloud-based infrastructure, platform, and software services, enabling scalability and interoperability.